


# Prompt Drift as a Reliability Risk in Software Engineering


An empirical investigation into  
the stability of LLMs in code  
generation and testing pipelines.

Ethan Chen  
Khoi Nguyen  
Pei-Yu Lin  
Pablo Rodriguez  
Shuang Ma

# The Problem



**15%**  
Accuracy  
Variance



Software Engineering relies on deterministic inputs and outputs. However, LLMs exhibit "Prompt Drift" - output instability even under rigid constraints. Performance gaps between best and worst runs can reach 70%.

Project Goat: Empirically measure drift across Code Gent Summarization, and Bug Detection.

# Research Questions

**1**

## Consistency

How consistent are code LLMs across repeated identical prompts?

**2**

## Sensitivity

How sensitive are models to semantically equivalent prompt paraphrasing?

**3**

## Model Comparison

Which models (GPT-4 vs. Llama-3/CodeLlama) exhibit the highest stability?


**4**

## Drivers

What task characteristics contribute most to instability?

# Paper 1 - Overview

Source: Atil et al., Non-Determinism of 'Deterministic' LLM Settings



Even with determinism is an engineering illusion.  
Instability is a system feature.

## Key Findings

1. Investigation: 5 Models x 8 Tasks x 10 Runs.
2. The Gap: Up to 15% accuracy variance across identical runs.
3. The Extremes: Performance gap between 'Best' and 'Worst' run reached 70%.
4. Conclusion: Non-determinism is likely inherent to efficient compute infrastructure (batching, hardware).

# Paper 1 - Metrics

**TARR** (Raw Agreement)

Output 1: “The value is 5”

Output 1: “The value is 5”

MATCH

**TARA** (Answer Agreement)

Output 1: “Answer: 5”

Output 1: “The result is 5”

SEMANTIC MATCH

We separate ‘string stability’ from ‘semantic stability’.  
A model can be accurate but unstable

# Paper 1 - Relevance

## System Factors

- Sparse Operations
- Continuous Batching

Atil et al. (2024)

Confirms ‘Drift’ is an infrastructure reality  
Establishes TAR baselines


## Inference Noise

## Methodology Control

- N=10 repeated runs per prompt
- Measure drift, not just accuracy
- Use TAR metrics for summarization

# Paper 2 Problem Definition

Errica et al., What Did I Do Wrong? Quantifying LLMs' Sensitivity and Consistency to Prompt Engineering, NEC.



**Insight:**

A tiny change can dramatically flip the model's prediction.

Accuracy alone does not tell developers how unstable a model is or why it fails.

**How can we quantify the sensitivity of an LLM to variations of the prompt?**

**Sensitivity**, measuring how much predictions change under semantically equivalent prompt variations

**Consistency**, measuring how stable predictions are across samples of the same class.

A slight change in the definition of the class "ENTY" causes a minor prompt variation that disrupts the Llama's prediction. This happens under the hood, making it very hard for a developer to debug the program.

# Paper 2: Methodology

## Metric 1: Sensitivity

How much predictions vary under prompt rephrasings

- Defined as **normalized entropy** of distribution:

$$S_\tau(\mathbf{x}) = -\mathbb{E}_{y \sim p_\tau(\cdot|\mathbf{x})} [\ln p_\tau(y|\mathbf{x})] / \ln(C),$$

- Expected sensitivity:

$$S_\tau = \mathbb{E}_{\mathbf{x}}[S_\tau(\mathbf{x})] \approx \frac{1}{N} \sum_{i=1}^N S_\tau(\mathbf{x}_i)$$

## Metric 2: Consistency

How similar predictions are across samples of the same class

- Total Variation Distance (TVD):

$$\text{TVD}(p, q) = \frac{1}{2} \sum_{c=1}^C |p(c) - q(c)|,$$


- Pair-wise consistency

$$C_y(\mathbf{x}, \mathbf{x}') = 1 - \text{TVD}(p_\tau(\cdot|\mathbf{x}), p_\tau(\cdot|\mathbf{x}'))$$


- Expected consistency:

$$C_y = \mathbb{E}[C_y(\mathbf{x}, \mathbf{x}')] \approx \sum_{\mathbf{x}, \mathbf{x}' \in \mathcal{D}_y} \frac{C_y(\mathbf{x}, \mathbf{x}')}{|\mathcal{D}_y|^2}.$$


# Paper 2: Results



(a) TREC




(b) CB




(c) DBpedia


Sensitivity for each sample of the dataset according to different prompting strategies.



(a) TREC



(b) DBpedia




(c) WoS

Violin plot of the Llama3 consistency over samples of the same classes, arranged by prompting technique, on different datasets.

## Paper 3 Motivation

### Measuring LLM Code Generation Stability via Structural Entropy

- LLMs often generate different code for the same prompt.
- This variability reduces reliability and reproducibility.
- Existing metrics focus on correctness, not consistency.
- They study how to directly measure code generation stability



# Paper 3 Methodology

## Phase1: Subtree Extraction

Input: Generation Code(Ta,Tb)

Action: Parse to AST -> Extract depth-bounded subtrees

Output: canonical encoding

## Phase2: Empirical Distributions

Actions: Count how often each subtree pattern appears.

Calculation: Count frequencies -> Normalize to vectors(P,Q)

Output: subtree frequency vector

## Phase3: Similarity Matrix

Action: Compare subtree distributions between code samples.

Metrics: Structural Cross-Entropy (SCE) and Jensen Shannon Divergence (JSD)

Result: Stability Score[0,1]

# Paper 3 Conclusion

Model	Language(Task)	Avg. BLEU	Code BLEU	Pass@1	Pass@5	TSED	SCE(structural)	SCE	JSD(structural)	JSD
LLaMA-3.1 8B it	Python	0.428	0.669	0.289	0.481	0.785	0.798	0.656	0.940	0.898
Qwen-2.5 7B it	BigCodeBench	0.596	0.716	0.339	0.482	0.765	0.823	0.726	0.947	0.913
Qwen-2.5-Coder 7B		0.614	0.715	0.373	0.517	0.764	0.823	0.722	0.951	0.918
LLaMA-3.1 8B it	SQL	0.495	N/A	0.673	0.824	0.832	0.729	0.669	0.934	0.905
Qwen-2.5 7B it		0.770	N/A	0.716	0.790	0.921	0.886	0.850	0.971	0.958
Qwen-2.5-Coder 7B	Spider	0.664	N/A	0.787	0.860	0.927	0.822	0.781	0.962	0.946

- Results show that even when pass@k is high, stability can still be low.
- JSD scores are generally high, meaning overall structure is similar.
- But SCE scores are lower, showing that token-level and fine-grained differences are common.
- This reveals instability that pass@k and BLEU miss.

# Paper 4: Motivation and Problem


ProSA: Assessing and Understanding the Prompt Sensitivity of LLMs

Main idea

- LLM answers change a lot when prompts are worded differently
- Current evaluations hide this problem

What the paper does

- Studies prompt sensitivity
- Looks at same question, different prompts
- Proposes a new metric: PromptSensiScore (PSS)



# Paper 4 Metrics and results

Key idea: Instance-level evaluation


- Not dataset averages
- Look at each question individually

PromptSensiScore (PSS)


- Compare model answers across multiple prompt versions
- High PSS = unstable model
- Low PSS = robust model

Main findings

- Larger models are usually more stable
- Math and reasoning tasks are very sensitive
- Adding few-shot examples helps a lot



# Paper 4: Explanation



## Why prompt sensitivity happens

- It's linked to model confidence
- Low confidence → answers flip easily

## Why this matters

- Benchmarks may overestimate model quality
- Users may get unreliable answers
- Evaluation should include prompt robustness

## Takeaway

- Prompt sensitivity is real
- It's measurable
- It should be part of LLM evaluation

# Paper 5: Motivation and Problem

## Benchmarks and Metrics for Evaluations of Code Generation: A Critical Review

TABLE I  
PERFORMANCE COMPARISON OF LANGUAGE MODELS

Model	Base Model	Size	Year	Benchmark	Score
GPT-NEO [18]	GPT-2	125M 1.3B 2.7B	2021	HumanEval	02.97 16.30 21.37
GPT-J [19]	GPT-2	6B	2021	HumanEval	27.74
Codex [20]	GPT-3	300M 679M 2.5B 12B	2021	HumanEval	36.27 40.95 59.50 72.31
TabNine	?	?	2021	HumanEval	7.59
ChatGPT	GPT-3.5	?	2022	HumanEval	94.00
Gemini-Ultra [17]	Transformer	?	2023	HumanEval Natural2Code	74.40 74.90
Gemini-Pro [17]	Transformer	?	2023	HumanEval Natural2Code	67.70 69.60
CodeGen [21]	Auto-regressive Transformer	350M 2.7B 6.1B 16.1B	2023	HumanEval	35.19 57.01 65.82 75.00
SantaCoder [22]	Decoder	1.1B	2023	MultiPL-E	45.90
InCoder [23]	Transformer	1.1B 6.7B	2023	HumanEval	25.20 45.00
StarCoder [12]	StarCoder-Base	15.5B	2023	HumanEval MBPP	33.60 52.70

## LLMs in Code Generations:

- Programming tasks categorized into D2C (Description to Code), C2D (Code to Description), and C2C (Code to Code). Code generation falls into D2C
- General purpose LLMs (ChatGPT) performs “better” in the benchmark, which conflicts with our assumptions
- HumanEval scores increased dramatically (2.97% to 94%) in one year. This signal potential problem in the accuracy of these benchmarks
- Model size and benchmark scores vary wildly. This is questionable comparability
- Benchmarks may not reflect *real* code usability or significance of differences

# Paper 5 - Methodology

How current benchmark work:

TABLE IV  
STRUCTURE OF DATA FOR EACH BENCHMARK

Benchmark	Context Code	#Test Cases	Solution
APPS	-	+	+
HumanEval	function signature	7.7	-
MBPP	-	3	-
MathQA-Python	-	3	-
ClassEval	class skeleton	33.1	+
CoderEval	function signature	+	-
Multipl-E	function signature	3 to 7	-
DS-1000	signature	1.6	+
HumanEval+	function signature	774.8	-
CONCODE	function signature	-	-
R-benchmark	-	-	+
JuICe	-	-	+
Exec-CSN	function signature	+	+
EvoCodeBench	function signature	+	+

- Data is sourced from code repositories, online forums, coding challenge sites, and textbooks.
- Tasks include a natural-language description and optional context code
- Context code define hint for the model and #test cases define evaluation rigor
- pass@k metric to estimate the probability of generating at least one correct solution in k trials

# Paper 5 - Conclusion

TABLE V  
EVALUATION PER BENCHMARK

Benchmark	Correctness	Metric	Model	Result
APPS	pass all tests	%pass@100	GPT-2 1.5B	0.68
			GPT-Neo 2.7B	1.12
HumanEval	pass all tests	AvgTPR	GPT-3 175B	0.06
			GPT-2 1.5B	7.96
MBPP	pass all tests	%pass@1	GPT-Neo 2.7B	10.15
			GPT-3 175B	0.55
MathQA -Python	pass all tests	%pass@1	GPT-Neo-2.7B	21.37
			GPT-J-6B	27.74
ClassEval	pass all tests	%pass@5	Tabnine	7.59
			Codex-12B	72.31
CoderEval (Python)	pass all tests	%pass@10	Decoder only-8B	79.0
			Transformer-68B	82.8
CoderEval (Java)	pass all tests	%pass@10	Lang. Model-137B	83.8
			Decoder only-8B	74.7
Multipl-E (HumanEval)	pass all tests	%pass@1	Transformer-68B	79.5
			Lang. Model-137B	81.2
Multipl-E (MBPP)	pass all tests	%pass@1	GPT-4	42.0
			GPT-3.5	36.0
DS-1000	pass all tests	%pass@1	WizardCoder	23.0
			StarCoder	14.0
HumanEval+	pass all tests	%pass@100	SantaCoder	10.0
			CodeGen	13.0
ConCode	BLEU	Avg BLEU	CodeGeeX	10.0
			InCoder	8.0
R-benchmark	Satisfactory	Avg #attempt <sub>k</sub>	Vicuna	4.0
			ChatGLM	3.0
ConCode	User-designed	Avg #attempt <sub>k</sub>	PolyCoder	3.0
			ChatGPT	23.48
HumanEval+	User-designed	Avg #attempt <sub>k</sub>	PanGu-Coder	27.39
			ChatGPT	30.00
ConCode	User-designed	Avg #attempt <sub>k</sub>	CodeGen	33.48
			PanGu-Coder	43.04
HumanEval+	User-designed	Avg #attempt <sub>k</sub>	ChatGPT	46.09
			CodeGen	≈ 36
ConCode	User-designed	Avg #attempt <sub>k</sub>	Codex	≈ 9
			InCoder	≈ 6
HumanEval+	User-designed	Avg #attempt <sub>k</sub>	CodeGen	≈ 40
			InCoder	≈ 14
ConCode	User-designed	Avg #attempt <sub>k</sub>	CodeGen	≈ 15
			InCoder	7.45
HumanEval+	User-designed	Avg #attempt <sub>k</sub>	CodeGen-002	41.25
			CodeGen-6B	8.4
ConCode	User-designed	Avg #attempt <sub>k</sub>	InCoder-6B	7.45
			CodeGen	≈ 64.0
HumanEval+	User-designed	Avg #attempt <sub>k</sub>	SantaCoder	40.6
			InCoder	≈ 29.8
ConCode	User-designed	Avg #attempt <sub>k</sub>	PolyCoder	13.6
			ChatGPT	89.8
HumanEval+	User-designed	Avg #attempt <sub>k</sub>	Vicuna	40.25
			StableLM- $\alpha$	11.9
ConCode	User-designed	Avg #attempt <sub>k</sub>	GPT-J	25.9
			GPT-Neo	16.8
HumanEval+	User-designed	Avg #attempt <sub>k</sub>	ChatGPT	1.6
			Retrieval	20.27
ConCode	User-designed	Avg #attempt <sub>k</sub>	Seq2Seq	23.51
			Seq2Prod	21.29
HumanEval+	User-designed	Avg #attempt <sub>k</sub>	User-designed	22.11

Current benchmark inadequacy:

- Model rankings change when testing rigor increases
- Benchmark results are highly sensitive to dataset design and sourcing
- The standard pass@k metric fails to reflect real-world usability, most don't re-run multiple times to find one solution
- The #attempt<sub>k</sub> metric better reflects productivity and whether code is a "useful starting point"

# Takeaway

**Drift Exists** (Atil et al.)

**Prompts Drive It** (Errica et al.)

**We Can Measure It** (Song et al.)

**Phrasing can change result**

**New metrics is required** (Paul et al.)

To transition LLMs from experiments to production components, we must treat '**Stability**' as a first-class metric alongside '**Accuracy**'.

**End of slides.**

# Our Experimental Design

	Models (tentative)	Tasks	Datasets
1	GPT-4 (Baseline)	Code Generation	HumanEval
2	Llama-3	Summarization	CodeSearchNet
3	CodeLlama	Bug Detection Test Generation	

## Timeline

